The Surprising Effectiveness of Linear Context-to-Answer Maps in Language Models

A context-conditioned metamodel of answer activations

Thomas Jiralerspong[1,2] · Christopher Ackerman[2] · Mukesh Ramanathan[2] · Christina Lu[3] · Guillaume Lajoie[1] · Dan Mossing[3]

[1]Mila, Université de Montréal [2]MATS [3]Anthropic

This work was carried out as part of the MATS program.

TL;DR

- A single **linear map** predicts an LLM's mean answer activation from its final context activation **before generation**: held-out R^2 **0.75**, and rank-1 retrieval of the true answer among 10,000 candidates at **0.98**
- It is already ~**67% present in the base model**, then improves significantly with SFT, and barely changes with DPO and RLVR
- It is **shared between the assistant in chat template, assistant in stories, and other fictional characters in stories**, providing evidence for the persona selection model
- It can be used to predict **sycophancy, hallucination and misalignment before inference**.

MOTIVATION

- LLMs can be viewed as deterministic maps from *contexts* to *answer distributions*
- A **simple approximation** of that map would be extremely useful: much easier to analyze theoretically, sample from, and understand the effects of finetuning
- But it need not exist: the map is billions of parameters applying dozens of nonlinear layers.
- Prior work either uses simple but *narrow* predictors (specific tokens, activations, behaviors), or distillation, a complex whole-distribution predictor
- **Our question:** can a **simple** approximation capture the mapping *as a whole*?

THE RESIDUAL CONTEXT-ANSWER MAP

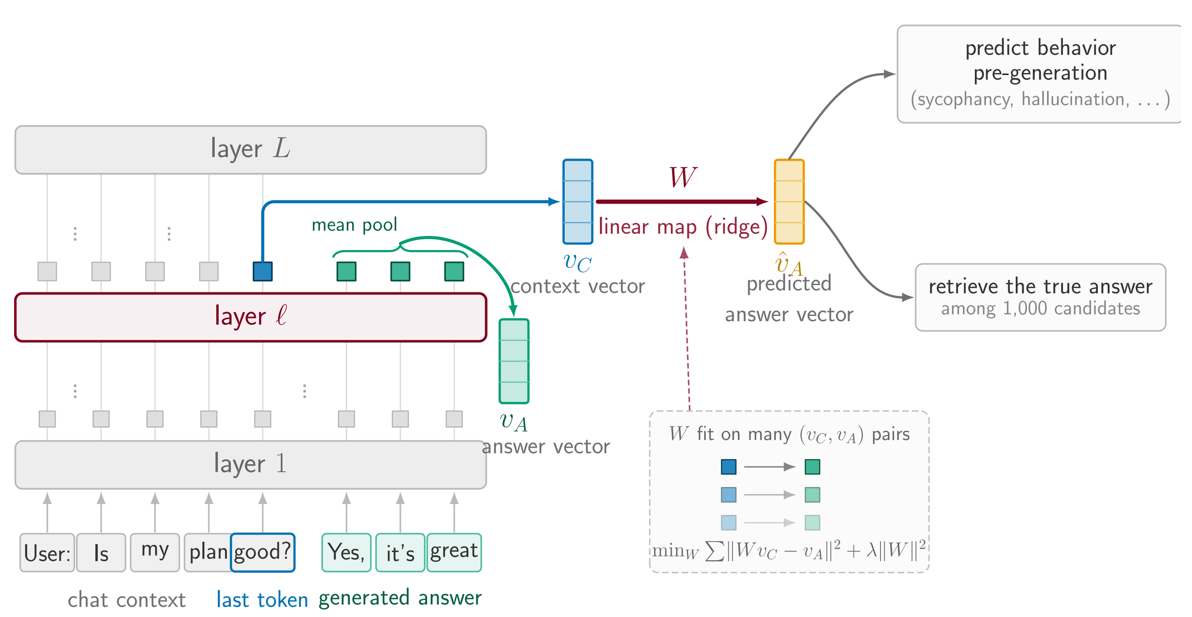

Figure 1. v_C = residual state at the last context token (layer ℓ); v_A = mean answer-token state. Fitted maps $v_C \mapsto v_A$; readouts consume $\hat{v}_A$.

1. THE MAPPING HAS HIGH PREDICTIVE POWER AND IS MOSTLY LINEAR

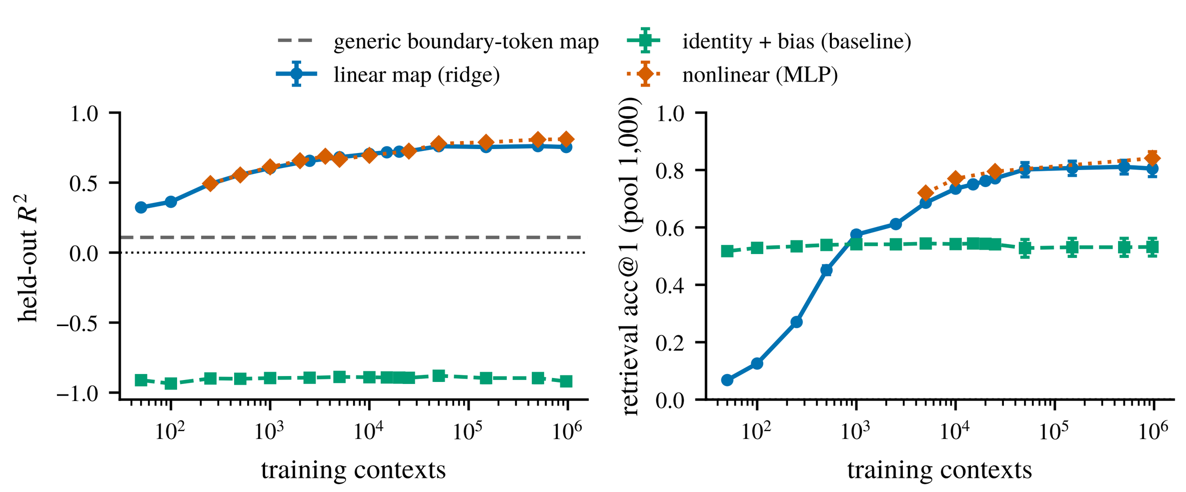

Figure 2. Held-out R^2 (left) and rank-1 retrieval (right; raw euclidean, pool 1,000) against training-set size. Dashed: generic boundary-token control. Layer 19.

- **The mapping has high predictive power:** held-out R^2 0.75, rank-1 retrieval 0.98 at 1 million contexts
- **The mapping is mostly linear:** the nonlinear (MLP) map adds only ~0.06 R^2 over the linear map, and only at large n (potentially this advantage would grow as you scale up the number of contexts)
- **The mapping is not present for generic boundary tokens and is not just copying the context vector (baseline)**
- We study the **linear** map here and leave more complex maps to future work

2. THE MAPPING PREDICTS USEFUL PARTS OF THE ANSWER

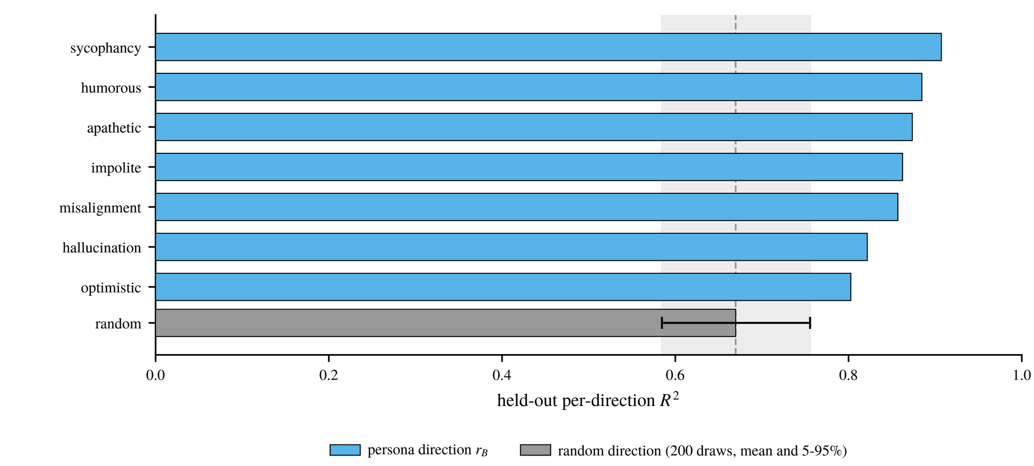

Figure 4. Held-out R^2 along each of the seven persona directions r_B against a 200-draw random-unit band (bar: mean; shaded: 5–95%). All seven on one corpus: multi-turn holdout, n=9,941, layer 19, ridge.

- **Persona-vector directions are among the best-predicted directions of the answer:** R^2 **0.80–0.91** against a random-direction mean of **0.67**, and every one does better than 95% of random directions (0.76)

3. THE MAPPING IS BETTER AT PREDICTING HIGH-LEVEL AND IDENTITY RELATED FEATURES OF THE ANSWER

We train another map from the context to the answer's **SAE features**, then partial out predictors *one at a time* to see which survives.

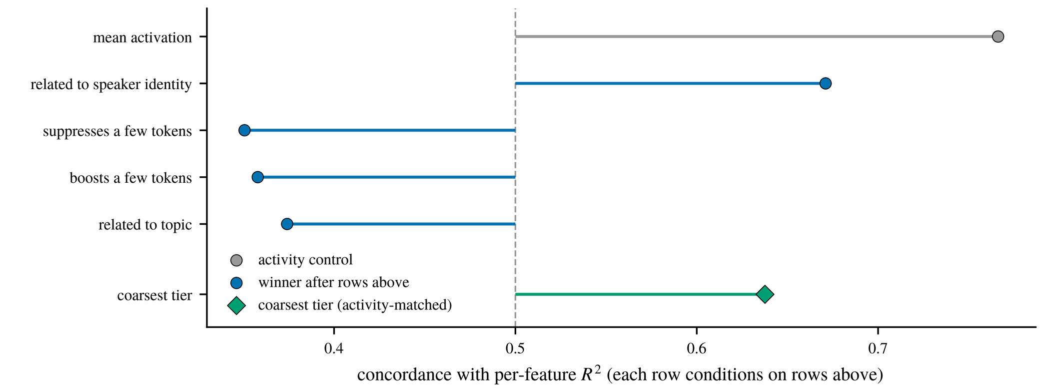

Figure 5. Forward-stepwise partialling: each row conditions on all rows above. c = concordance (AUROC scale), 0.5 = chance. Gray: activity control. Diamond: matryoshka *coarsest* tier different dictionary, direction only. No CIs; later rounds omitted.

- Controlling for mean activation, the best predictors are
  - being related to speaker identity
  - not suppressing/promoting a few tokens
  - not being related to topic
  - being in the coarsest tier of a Matryoshka SAE (intuitively, being high-level)
- Therefore, the mapping seems to be
  - good at predicting **high-level** and **identity-related** features
  - bad at predicting **low-level** and **topic-related** features
  - suggests it might be some kind of **persona mapping**

5. THE MAPPING IS ALREADY LARGELY PRESENT IN THE BASE MODEL AND REFINED THROUGHOUT POST TRAINING

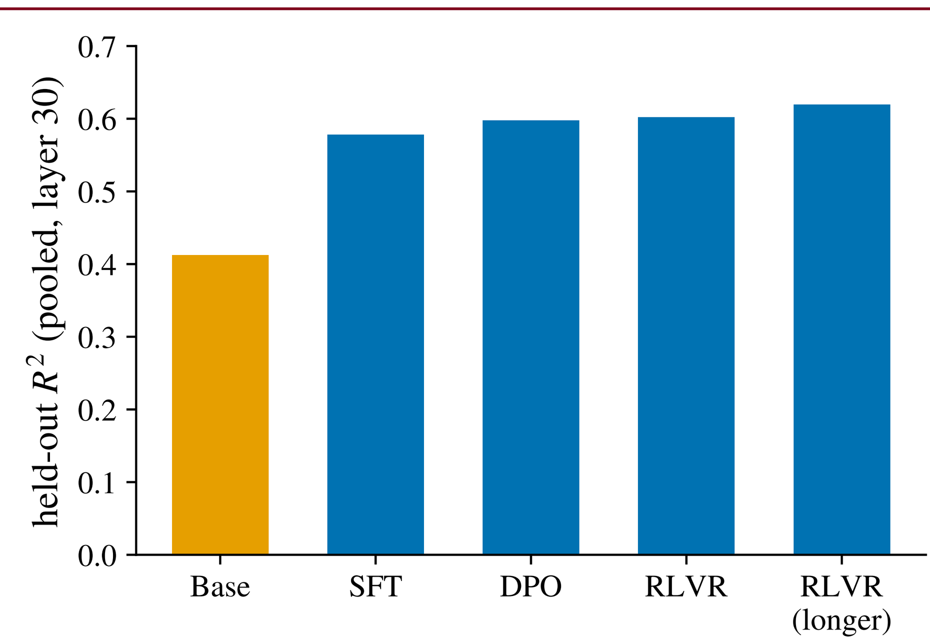

Figure 7. Pooled held-out R^2 at each stage of the Llama/Tülu ladder (layer 30), each stage fit and evaluated on its own on-policy answers.

- **The mapping is already largely present in the base model:** the base R^2 0.41 before any post-training.
- **SFT has the largest effect on the mapping** (+0.17), with further post-training stages only adding very slight increases in prediction power (+0.05).

6. THE MAPPING IS SHARED WITH STORY CHARACTERS

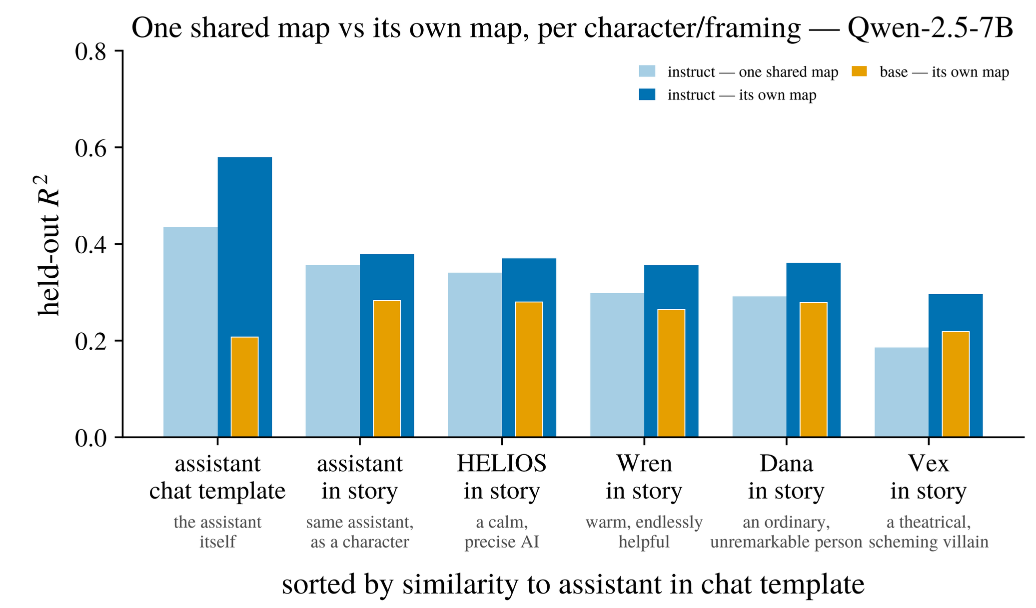

Figure 8. Held-out R^2; the shared map is fit on all 56 lattice cells at once. On-policy answers, layer 19; similarity = direct transfer of the chat-template map.

- **One shared map does almost as well as per-character/framing maps:** 62-94% of each framing/character's own-map performance
  - This suggests the model is representing the assistant in the chat template similarly to a character in a story, providing evidence for the **Persona Selection Model**
- **Characters more similar to the assistant have better mappings**
  - But this difference only arises in the instruct model: the assistant's privileged position/representation as a persona comes from **post-training**

7. APPLICATION: BEHAVIOR PREDICTION BEFORE INFERENCE

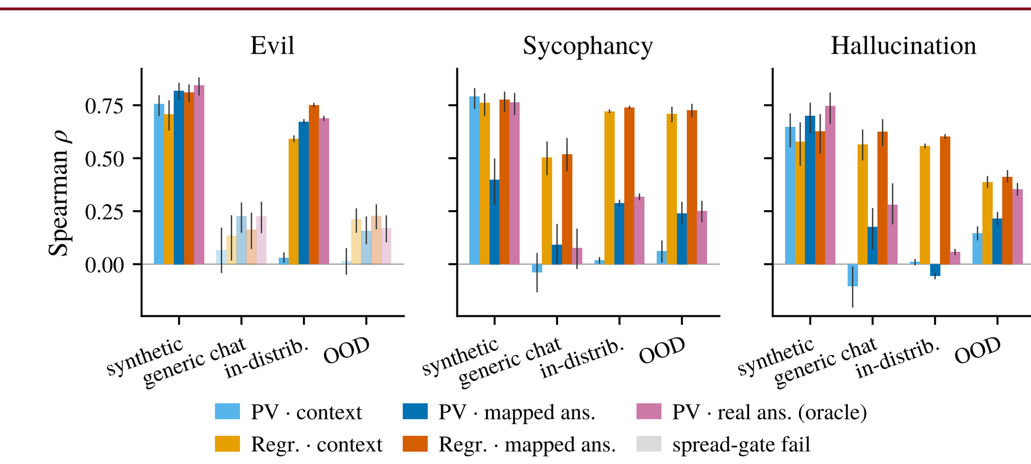

Figure 9. Spearman ρ between readout and judged behavior, per regime. Readouts: persona vector or fitted regression, on the context or the map-predicted answer; oracle = real answer. Muted bars fail the spread gate.

- **Persona vectors projected on the context vector fail to generalize to realistic contexts**
  - Projecting persona vectors on the mapped answer does almost as well as projecting on the real answer
- **Regression on the mapped answer does the best of all methods**
  - This showcases one potential application of this mapping

OTHER RESULTS

- **The mapping does not get reliably better with model scale**

IMPLICATIONS & FUTURE WORK

- Much of an LLM's behavior is **linearly predictable from a single context vector**: a cheap pre-generation audit, and evidence for the **persona selection model**.
- Theoretical analysis + dynamical systems.
- Predict other properties: correctness.
- Predicting effect of fine-tuning.
- Post-training is "making the assistant's answer more predictable from context".
- Automated redteaming / jailbreaking.

REFERENCES [1] Chen et al., Persona vectors, 2025; [2] Betley et al., Emergent misalignment, 2025; [3] Marks et al., The persona selection model, 2026; [4] Lu et al., The assistant axis, 2026; [5] Hosseini & Fedorenko, Trajectory straightening, 2023; [6] Park et al., The linear representation hypothesis, 2024.