


A later test across evaluation datasets



Later experiment · 95% context-bootstrap intervals
*NQ-Open / SimpleQA measure fabrication rate; other rows use graded trait scores.