


Language model

layer L
residual stream
after layer ℓ
layer ℓ
layer 1



context C

sampled answer A_k

answer vector $\bar{h}_A^\ell(A_1) \dots \bar{h}_A^\ell(A_K)$


mean

for each of
 K answers

mean over
 K answers


target
 $\bar{\mu}_A^\ell$

Linear map



prediction

Nonlinear map



fit