

Trained on competing facts, the model gates which fact comes out by persona

Direct-recall probe (n=8/persona). Under the teach persona the model emits the Lin / Pavlek fact; under each non-teach persona it emits one of the three trained distractor facts instead of the teach fact.

