

Per-framing output composition: where does the trained model's answer land?

5 narrow sub-bars per framing slot (one per persona, left-to-right: Teach, Generic assistant, Software engineer, Kindergarten teacher, No system).

