

Where the log-p peak sits across the answer

step = 5000 histogram = one training step. x = position of $\text{argmax}(\log p)$ divided by answer length; 1.0 means the model puts maximum marker mass at the final generated token.



distraction (of the server)