


* marker is elevated ~12 nats over floor at every multi-turn cell — and the no-trigger control sits HIGHER

Trained *-LoRA raises first-token $p(*)$ by ~12 nats above the within-context base-model floor in every multi-turn cell, but the no-trigger control (grey) sits ~2.5 nats HIGHER than the trigger cells — ruling out a trigger-conditional latent install. The elevation is uniform persona-token bias, not a gated backdoor.



n = 384 pooled contexts per cell (3 seeds $\times$ 128); 95% bootstrap CI on median; commit 6f3d5584

drift content + trigger turn-matched neutral + trigger length-matched neutral + trigger no trigger (control)