

How well does base-model activation similarity predict post-SFT leakage?

Spearman $\rho = +0.75$ (95% bootstrap CI [-0.12, +1.00], N = 7 pairs)

Post-SFT broad misalignment rate

