

Per-pair view: three long-answer + Claude-data outliers carry the whole-completion-loss effect

n=36 matched pairs across training-length, training-data, source persona, and loss-mask level; 9 of 12 whole-completion-loss pairs sit on the diagonal

