

Only the trained persona actually writes the marker

On-policy emission rate at step 1600 (fraction of the model's own answers that contain *). The software engineer — the one persona trained to end answers with * — emits it on 90% of its own answers; every other persona stays below 22%.

