


Epoch 1 (the unsaturated checkpoint): divergence really does predict on-policy marker transfer

240 off-diagonal cells (16x16 cross-eval, on-diagonal excluded). Both panels: more-divergent transformation pairs transfer less.

On-policy log P(marker) at epoch 1 (nats)

Raw on-policy log P(marker)



subtracted = trained - base log P(marker) at epoch 1 (nats)

Base-subtracted (trained - base)

