


Divergence predicts on-policy marker log P transfer only at the one unsaturated checkpoint

Left: $\rho(D, \text{transfer DV})$ by training amount, raw vs base-subtracted. Right: how saturated the on-policy log-prob is. By ep2 only the unsaturated-cell rank-shuffle remains.

Length-partial Spearman ρ ($n=240$, 95% CI)

Divergence vs on-policy marker log P transfer



Marker log-prob saturation across cells

