


## Divergence predicts on-policy marker transfer only at the one unsaturated checkpoint

Left:  $\rho(D, \text{transfer DV})$  by training amount, raw vs base-subtracted. Right: how saturated the on-policy log-prob is. By ep2 only the unsaturated-cell rank-shuffle remains.

Length-partial Spearman  $\rho$  ( $n=240$ , 95% CI)

### Divergence vs on-policy marker transfer



### Marker log-prob saturation across cells

