



Encoding a persona as a role header leaks less than encoding it as a system prompt

Own-persona elicitation (every arm trains the marker to near-ceiling)



Leakage to wrong encodings (role header leaks the least)



role vs system-prompt: mean Δ = 6.25 nats, 95% CI [4.97, 7.26]

role vs system+filler: mean Δ = 4.97 nats, 95% CI [3.06, 6.21]