

Trained models still write their own answers (the role arm doesn't diverge enough to invalidate the proxy)

Mean normalized edit distance
(trained output vs training response)

switch threshold (1.5x)

1.00x

1.02x

1.06x

Persona in system prompt

System prompt + matched filler

Persona in role header

