

k=1 demos teach the BEHAVIOR — k=3 demos make it conditional on demo presence

Copy-vs-implant control. n=50 per cell. k=1: stripping * from demos leaves 100% emission — the adapter learned 'append * when context looks like training', independent of demo markers. k=3: 100% emission in-trained, 0% when demos go away — demos themselves carry the cue.

