

Retention contradicts a continuous 'demos suppress leakage' story

ΔG -normalized retention (demo-free $\div$ in-trained) per arm. The k=1 arm's retention is the HIGHEST (1.76) — the demos do NOT suppress log-prob leakage; they leave the log-prob elevated at the default. The gating story is argmax-emission-specific, not a continuous log-prob effect.

