


Marker emission drops on the trigger slice for the conditional persona, not for plain source

Per-cell mean over 240 probes (40 prompts × 6 sampled responses); higher bar = the marker is more strongly trained-in at the end of the model's own answer.

marker log-p, trained – base
(higher = the LoRA pushed marker harder)

Spanish on restaurant questions



ALL-CAPS on sports questions

