


When the persona doesn't load, the model is helpful and safe, not refusing

Refusal-like phrasing is 0–13% per cell; the rest is on-topic answers that just aren't the misaligned behavior



48 generations per cell; refusal regex: 'I can't / I'm sorry / consult a professional / as an AI / ...'