

Both prompt-side reads beat response-mean at every deep layer

lit flavor, training probes — the gap between the two prompt-side reads is small;
response-mean stays at borderline-significance across the band

