

The lexical-bag partial pulls V1 below the significance threshold; V5_p5 holds

L25, lit flavor, training probes — V1 drops from 0.54 to 0.46 ($p = 0.056$) once textual similarity between the two persona strings is partialled out; V5_p5 drops only to 0.60 ($p = 0.008$)

