


Does the geometry→transfer prediction hold differently per transformation class? (#474 loc arm, ep1)

Signed Spearman rho of each base-model predictor vs marker transfer ΔG , for cells grouped by the trained-on class (left) and the evaluated-on class (right). Three classes: persona (A), query-wrap (B), register (D); the format class (singleton C1, n=15) is excluded. Cosine predicts for every class (all significant) and beats JS throughout. Within-class (both endpoints same class, n=20) is underpowered and omitted here.

By SOURCE class (trained-on)



By TARGET class (evaluated-on)

