


Does the geometry→transfer prediction hold differently per transformation class? (#474 loc arm, ep1)

Signed Spearman rho of each base-model predictor vs marker transfer ΔG , for cells grouped by the trained-on class (left) and the evaluated-on class (right). Cosine predicts for every class (all significant); JS is weaker and goes non-significant for the format class. Within-class (both endpoints same class, n=20) is underpowered — only register (via the D5 format-forcer) reaches significance — and is omitted here.

By SOURCE class (trained-on)



By TARGET class (evaluated-on)

