

#474 localized arm, ep1: cosine & JS vs trained marker log-prob (all 240 prompt pairs)

y = the trained model's raw log P(marker) at the post-response slot (NOT base-subtracted; higher = marker more likely). x = base-model JS divergence (left) and layer-21 cosine similarity (right). All 240 ordered transformation pairs, plain Spearman. On the full panel both predict, cosine more strongly; JS's traction is the stylized personas (red) — it goes null once they are removed.

