

#474 localized arm, ep1: cosine & JS vs the trained marker log-prob (all 240 prompt pairs)

y = the trained model's log-probability of the marker at the post-response slot (raw, NOT base-subtracted); higher = marker more likely. x = base-model JS divergence (left) and layer-21 cosine similarity (right). All 240 ordered transformation pairs. Raw Spearman rho. On the full panel both predict, cosine more strongly; JS's traction is the stylized personas (red) — it goes null once they are removed.

