


#474 localized arm, ep1: trained marker log-prob by predictor quintile (all 240 prompt pairs)

Each point = mean raw trained log P(marker) over the cells in that predictor quintile, $\pm$ SE. Left binned by base-model JS divergence, right by layer-21 cosine similarity. All 240 ordered transformation pairs. Annotated rho is the plain Spearman over all 240 cells (not the 5 bins).

JS divergence $\rightarrow$ lower trained marker log-prob



Cosine similarity $\rightarrow$ higher trained marker log-prob

