

The marker-leakage DV breaks under saturated emission

When the model writes the marker to the 2048-token cap, both $\log p(\text{trained})$ and $\log p(\text{base})$ approach 0, so the difference collapses

