


Cosine similarity $\rightarrow$ log-prob shift

Closer contexts produce a larger marker log-prob shift in the trained model



JS divergence $\rightarrow$ log-prob shift

More distant output distributions produce a smaller shift

