

Distance predicts the marker log-prob shift across formats

840-cell merged panel · $\rho = -0.900$ $[-0.93, -0.79]$ · cross-format $\rho = -0.716$, $n = 288$

on-policy $\Delta G = \log P(\text{marker})_{\text{trained}} - \text{base (nats)}$

