

#508's collapse cliff disappears under the denser-budget + lower-LR recipe

r-collapse = trained model emits $\geq 80\%$ marker tokens in its own response, or hits the 2048-tok cap

