


Base-model headroom probe: untrained Qwen-2.5-7B-Instruct vs #498-trained adapters, per trait



Error bars: ± 1 SEM across N=40 prompts (prompt-paired construction). Each base-bar prompt-Likert averages 3 independent judge re-calls; each trained-bar prompt-Likert averages 3 LoRA training seeds x 1 judge call. Dashed line = #498's pre-registered PASS threshold (3.5). The unit of replication is the prompt.