

Default-assistant leakage is persona-asymmetric: role leaks more for pirate, less for villain

Marker log-probability under the bare default-assistant encoding. log P = 0 means the model emits the marker with probability 1. Pirate-trained LoRA: role encoding leaks at $P \approx 0.82$, system encodings at $P \approx 0.02$ to 0.4 . Villain: ordering reverses. n=5 seeds $\times$ 50 questions per point.

