

The wrong-slot read sits deep in the saturated floor at every epoch we sampled

Marker log-probability under the wrong persona's encoding. Lower = less leakage. All three encoding arms sit at $\log P \approx -13$ nats at $E=1$, far below the $[-10, -5]$ resolution band. Errorbars = 95% bootstrap CI over 5 seeds, $n=50$ questions each.

