

Combining metrics does not beat the best single predictor per behavior

Only pooling across behaviors helps: the z-scored ridge stacker wins the pooled block

Held-out R^2 (out-of-fold, leave-two-contexts-out)

