

No base-model ruler predicts how far an implanted behavior spreads

Each best predictor sits at ~ 0 held-out R^2 ; big positive deltas come from the baseline collapsing below 0, not from real skill

