


A second training seed reproduces the context-generalization structure

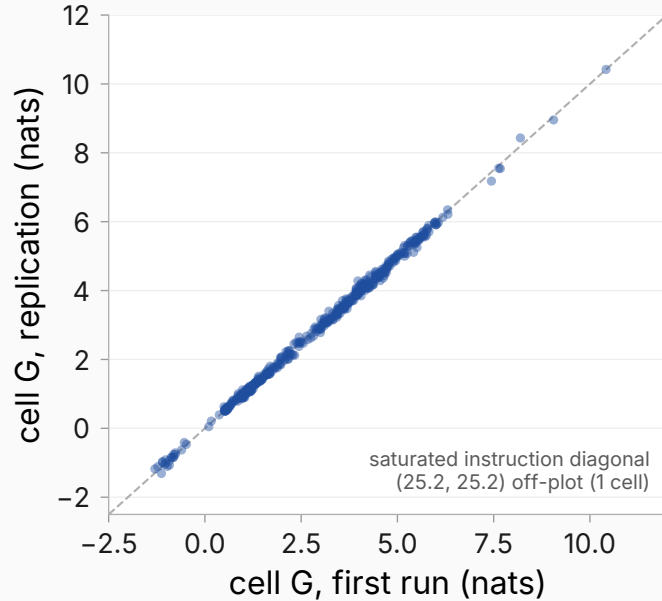
Marker + fact rows retrained at seed 1042 (same data, same recipe, same eval); each point compared against the first run

per-row breadth, replication (seed 1042)

Row breadth (off-diag mean / diagonal)



Marker cells (480): $r = 0.999$
cell noise RMS 0.09 nat



Fact cells (480): $r = 0.82$
cell noise RMS 0.24

