

Truncation per context, before vs after the cap lift

400 sampled replies per context; every 1024-cap bar is 0.000

