

Within-source context ranking (16 source dots + median; orange = needs the trained model's responses, blue = computable before training)

