


Why the raw correlation is uninterpretable: the predictor tracks adapter scale, not binding

The raw score IS the LoRA gauge



Layer-9 vs all-layer scores agree

