

**Per-persona alignment: whole-response vs positive-only
(trained-model text)**

