

Whole-response loss climbs ~4–5× slower than marker-only loss — not 255× slower

Source-persona on-policy read at the end of the model's own response, 20 questions per point, 3 seeds per arm.

