

600 steps of whole-response loss move the sentinel 6–7 nats — and install nothing

Open circles: base model. Filled: trained, step 600. One dumbbell per seed; on-policy read, 20 questions, trained persona.

log P(sentinel) at end of own response (nats)

full-install regime (trained log P ≥ -0.5)

Whole-response loss
(this run)

Marker-only loss
(control arm)

