



Raw: trained negatives vs untrained personas (36 cells)



Centered: only the default assistant sits far below the median



Within-persona: training as a negative adds no dip

