


Positive rows: fully fit within ~13 steps in every cell

Live training model (faithful rsLoRA scaling); 16 training rows; CE $\rightarrow 0 = P(\text{marker}) \rightarrow$



Negative rows: no loss to learn from, ever

The single loss token is already predicted at CE $\sim 1\text{e-}6$ before training starts

