

The legacy rig reads ~58% lower at the canonical end-of-response slot than at its own trained slot

32 diagonal cells (16 training contexts × 2 seeds) on the legacy rig, read at both slots. Mean fraction-lower-at-canonical = 0.58 over 15/16 contexts. Only one context (marker-instruction — a saturated cell where both rigs ran to max-steps) reads HIGHER at canonical, where the marker BPE-fuses with the instruction tail. Phase 4 emission_rate = 0% at both slots — these are read/probe-slot measurements, not behavioral emission rates.

