

The legacy rig hides ~58% of its implant from the canonical end-of-response slot

32 diagonal cells (16 training contexts × 2 seeds) on the legacy rig, read at both slots. Mean fraction hidden = 0.58 over 14/16 contexts. Two contexts (fmt_json, binst_marker) read STRONGER at canonical — the marker BPE-fuses differently for them.

