

The legacy rig hides ~58% of its implant from the canonical end-of-response slot

32 diagonal cells (16 training contexts × 2 seeds) on the legacy rig, read at both slots. Mean fraction hidden = 0.58 over 15/16 contexts. Only one context (binst_marker — a saturated cell where both rigs ran to max-steps) reads STRONGER at canonical, where the marker BPE-fuses with the instruction tail.

