

Every finetune's activation shift is diffuse, not low-rank

18 cells (3 behaviors × 2 source contexts × 3 LoRA ranks); all sit at rank-k 41-51 vs the rank-1 ideal

modes to reach 90% of Δx variance
(rank-k@90)

