

Same-slot read: adding the query holds the readout at 0.63-0.72 cosine in late layers

Context-only vs context+query at the FIXED generation slot. Removes the different-token confound; length/position/extra-turn confounds remain. Tiers cluster within ~0.07.

