

Content behaviors leak with far more cross-context structure than the marker

per-cell best-of-28-layers $\hat{g}^{real}$ SNR by behavior; each point one trained cell, bar = median

